POOJITHA MARIDI

Middletown, DE | +1 (302) 363-9177 | poojitham0914@gmail | LinkedIn | Portfolio

PROFESSIONAL SUMMARY

- Data Engineer with 7+ years of experience designing and delivering scalable ETL pipelines and cloud data platform workflows using Azure Data Factory, Azure Databricks, Azure Data Lake Storage, AWS cloud services, Spark SQL, and SQL Server supporting enterprise analytics environments.
- Developed metadata-driven ingestion pipelines and distributed Spark transformations improving pipeline reliability and reducing manual validation effort by 30% across Business Intelligence dataset layers.
- Optimized SQL Server and Hive-based processing workflows through indexing strategies and schema validation standardization reducing reporting latency by 34% across analytics delivery pipelines.
- Collaborated with Analytics Engineers and DataOps teams to deploy version-controlled pipelines using Azure DevOps while securing workflow credentials through Azure Key Vault and Logic Apps across multi-environment deployments.

TECHNICAL SKILLS

Cloud Data Platforms: Azure Data Factory, Azure Databricks, Azure Data Lake Storage, Azure Key Vault, Azure Logic Apps, Azure DevOps, AWS.

Data Engineering and DataOps: Data Engineering, Data Pipeline Engineering, Data Platform Engineering, ETL Pipeline Development, Metadata-driven Pipelines, Batch Processing Pipelines, Workflow Orchestration, Data Integration, Schema Validation, Data Quality Validation, Pipeline Monitoring, CI/CD for Data Pipelines, Pipeline Release Management, Dataset Curation Layers, Database Engineering.

Big Data Technologies: Apache Spark, Spark SQL, Hive, Apache Kafka, Distributed Data Processing, Partitioning and Bucketing Strategies.

Programming and Query Languages: Python, SQL, Java, PL/SQL.

Databases and Data Warehousing: SQL Server, Oracle Database, MySQL, MongoDB, Relational Database Design, Query Optimization, Structured Dataset Modeling.

Analytics and Business Intelligence Support: Business Intelligence Dataset Preparation, Analytics Engineering Support, Reporting Data Transformations, Grafana Dashboards.

Engineering Collaboration and Delivery Tools: Git, Maven, JIRA, Confluence.

PROFESSIONAL EXPERIENCE

Senior Data Engineer **Apr 2025 - Present**
Indus Group | Middletown, DE

- Architected scalable Azure Data Factory batch ingestion pipelines integrating structured database and Azure Data Lake Storage sources across hybrid Azure and AWS cloud environments increasing pipeline throughput capacity by 35%.
- Built distributed Spark-based transformation workflows in Azure Databricks preparing analytics-ready datasets accelerating Business Intelligence dashboard refresh cycles by 27%.
- Coordinated Kafka-supported ingestion sequencing improving near real-time dataset availability across cross-functional analytics consumption workflows.
- Optimized SQL Server transformation procedures through indexing strategies improving reporting query execution performance by 31%.
- Automated pipeline monitoring and alert orchestration using Azure Logic Apps enabling DataOps teams to reduce incident response turnaround time by 29%.
- Deployed version-controlled ETL releases through Azure DevOps CI/CD pipelines improving reliability across multi-environment Data Platform deployments by 25%.
- Standardized schema validation frameworks using Spark SQL and Hive table-layer integrity checks improving ingestion data quality consistency by 24%.
- Partnered with Database Engineers and Analytics Engineers to implement structured transformation layers supporting regulatory and operational reporting workloads reducing reconciliation effort by 26%.

Data Engineer **Feb 2024 - Mar 2025**
Raiden Tech Group | Middletown, DE

- Engineered Azure Data Factory ETL pipelines integrating SQL Server sources with Azure Data Lake Storage across hybrid Azure and AWS cloud environments improving ingestion reliability across analytics workflows by 32%.
- Developed Spark SQL transformation notebooks in Azure Databricks to standardize schema validation across ingestion layers reducing Business Intelligence dataset defects by 26%.
- Optimized relational database workloads using indexing and execution-plan tuning improving dashboard refresh performance across reporting environments by 28%.
- Implemented metadata-driven orchestration frameworks using Azure Data Factory ForEach and Execute Pipeline activities reducing pipeline maintenance effort by 22%.
- Secured pipeline credential access through Azure Key Vault integrations strengthening governance compliance across controlled Data Platform environments.
- Coordinated ingestion readiness with Analytics Engineers to prepare curated transformation layers improving downstream Business Intelligence dataset availability by 24%.
- Integrated Hive-based staging tables supporting distributed dataset structuring across batch processing layers improving ingestion-layer consistency across Big Data workflows.
- Automated pipeline execution tracking using Azure DevOps release workflows improving deployment traceability across ETL lifecycle environments.

Graduate Teaching Assistant **Aug 2022 - Dec 2023**
University of North Texas | Denton, TX

- Supported Big Data coursework by mentoring students on Spark SQL-based ETL pipeline assignments aligned with modern Data Engineer workflow design patterns, improving lab completion success rates by 22%.
- Guided learners through SQL-driven transformation exercises reinforcing database engineering concepts used in analytics engineering and Business Intelligence dataset preparation.
- Assisted faculty in preparing distributed datasets and evaluating Spark DataFrame submissions across Data Platform-oriented programming labs supporting hands-on pipeline experimentation.
- Coordinated Git-based project validation workflows with Software Engineering instructors to review schema integrity across student-built data pipeline implementations.

- Facilitated troubleshooting sessions for Python-based transformation scripts used in analytics processing environments, reducing assignment resubmission rates by 18%.
- Maintained Hive-enabled lab environments supporting distributed storage experimentation for Big Data processing and ETL workflow simulations across classroom clusters.

Data Engineer **Jun 2018 - Nov 2021**
Bank of America | India

- Designed Azure Data Factory ETL pipelines to migrate structured datasets from on-prem SQL Server into Azure Data Lake Storage supporting enterprise Data Platform modernization and reducing manual ingestion effort by 30%.
- Developed Databricks notebooks using Spark SQL and Python to automate rule-based validation across transaction datasets enabling Analytics Engineering teams to improve downstream Business Intelligence readiness.
- Optimized high-volume SQL Server database queries through indexing strategies and execution-plan tuning, decreasing compliance analytics reporting latency by 34%.
- Implemented metadata-driven Azure Data Factory workflows using ForEach, Filter, Stored Procedure, and Execute Pipeline activities to orchestrate scalable Data Pipeline Engineer ingestion patterns across regulated environments.
- Integrated Azure Key Vault with Azure Logic Apps to secure credential access and automate ETL pipeline failure notifications to DataOps support teams, improving incident response turnaround time by 25%.
- Standardized schema validation logic within Databricks processing layers by collaborating with database engineering stakeholders, improving consistency across multi-source data integration pipelines.
- Deployed version-controlled ETL releases through Azure DevOps enabling structured DataOps lifecycle management and reducing production deployment errors by 20%.
- Partnered with Business Intelligence stakeholders to transform curated Azure Data Lake datasets using Spark transformations, accelerating dashboard refresh availability for regulatory monitoring workflows.

Big Data Intern **Apr 2017 - Jul 2017**
Efftronics Systems Pvt. Ltd. | India

- Developed structured ETL workflows using Python and SQL to extract datasets from MySQL databases and prepare analytics-ready tables powering Grafana Business Intelligence dashboards.
- Implemented Spark SQL transformation strategies including partitioning and broadcast joins to optimize batch processing performance across distributed Big Data datasets by 28%.
- Engineered Hive external and managed tables supporting scalable storage layers for telemetry-driven Data Platform ingestion pipelines across analytics environments.
- Automated schema cleansing routines using Python validation scripts to standardize multi-source datasets prior to loading into reporting pipelines.
- Integrated Google Analytics streams into Grafana visualization layers enabling cross-platform Business Intelligence monitoring for operational teams.
- Collaborated with database engineering stakeholders to optimize PL/SQL transformation logic supporting structured SQL-based ETL execution across ingestion workflows.

PROJECTS

Enterprise Data Lake Ingestion Framework

- Designed a reusable ingestion framework in Azure Data Factory to consolidate multi-source SQL Server datasets into Azure Data Lake Storage improving onboarding speed for new datasets by 40%.
- Implemented parameterized pipeline templates using metadata-driven configurations enabling scalable ingestion across multiple schema variations without duplicating workflow logic.
- Coordinated dataset readiness validation with reporting stakeholders ensuring structured transformation layers supported downstream Business Intelligence consumption cycles.

Distributed Transaction Data Transformation Pipeline

- Developed Spark-based transformation workflows in Azure Databricks to standardize high-volume transactional datasets into analytics-ready formats supporting enterprise reporting pipelines.
- Applied Spark SQL partitioning strategies to restructure large datasets improving transformation execution efficiency across distributed compute clusters by 30%.
- Validated schema consistency across curated transformation layers enabling Analytics Engineers to accelerate dashboard dataset integration timelines.

Real-Time Event Streaming and Hive-Based Staging Pipeline

- Established Kafka ingestion coordination pipelines to support near real-time dataset availability across streaming-aware analytics environments.
- Engineered Hive staging tables to organize intermediate ingestion layers improving structured dataset accessibility across distributed processing workflows.
- Enabled cross-team dataset synchronization by integrating structured staging outputs with reporting delivery pipelines supporting operational monitoring use cases.

EDUCATION

Master of Science in Computer Science **Dec 2023**
University of North Texas | Denton, TX

Bachelor of Technology in Computer Science **May 2018**
VR Siddhartha Engineering College | India

CERTIFICATIONS

- **Microsoft** Certified: Azure Data Engineer Associate
- **AWS** Certified Solutions Architect Associate
- IBM Data Engineering Professional Certificate - **Coursera**
- Google Cloud Data Engineering Professional Certificate - **Coursera**
- Data Warehousing for Business Intelligence Specialization - **Coursera**
- Big Data Specialization - **Coursera**